

DeceSSIONTreeClassifier

Akshat Kumar (242211003) 12-02-2025

```
# Libraries
import pandas as pd
from time import time,ctime
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn import metrics
```

```
print("Timestamp: "+ctime(time()))
```

```
Timestamp: Wed Feb 12 06:27:22 2025
```

```
df = pd.read_csv('diabetes.csv')
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Next steps:

Generate code with df

View recommended plots

New interactive sheet

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null    int64
1   Glucose                768 non-null    int64
2   BloodPressure          768 non-null    int64
3   SkinThickness          768 non-null    int64
4   Insulin                768 non-null    int64
5   BMI                    768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                    768 non-null    int64
8   Outcome                768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

Data Exploration and Preprocessing:

```
df.shape
```

```
(768, 9)
```

```
df.columns
```

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

```
df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

```
df.Outcome.value_counts()*100/len(df)
```

	count
<b>Outcome</b>	
0	65.104167
1	34.895833

```
df.groupby('Outcome').mean()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
<b>Outcome</b>								
0	3.298000	109.980000	68.184000	19.664000	68.792000	30.304200	0.429734	31.190000
1	4.865672	141.257463	70.824627	22.164179	100.335821	35.142537	0.550500	37.067164

```
df.groupby('Outcome').agg(['mean','median'])
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPed
	mean	median	mean	median	mean	median	mean
<b>Outcome</b>							
0	3.298000	2.0	109.980000	107.0	68.184000	70.0	19.664000
1	4.865672	4.0	141.257463	140.0	70.824627	74.0	22.164179

## 2. Handle missing values (if any) appropriately.

```
df.isnull().sum()
```

	0
<b>Pregnancies</b>	0
<b>Glucose</b>	0
<b>BloodPressure</b>	0
<b>SkinThickness</b>	0
<b>Insulin</b>	0
<b>BMI</b>	0
<b>DiabetesPedigreeFunction</b>	0
<b>Age</b>	0
<b>Outcome</b>	0

## Model Implementation:

✓ 1. Split the dataset into training and testing sets (e.g., 80-20 split).

```
Y = df['Outcome']
Y
```

	Outcome
0	1
1	0
2	1
3	0
4	1
...	...
763	0
764	0
765	0
766	1
767	0

768 rows × 1 columns

```
X = df.drop('Outcome',axis=1)
X
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33
...	...	...	...	...	...	...	...	...
763	10	101	76	48	180	32.9	0.171	63
764	2	122	70	27	0	36.8	0.340	27
765	5	121	72	23	112	26.2	0.245	30
766	1	126	60	0	0	30.1	0.349	47
767	1	93	70	31	0	30.4	0.315	23

768 rows × 8 columns

Next steps: [Generate code with X](#) [View recommended plots](#) [New interactive sheet](#)

```
X_train , X_test, y_train, y_test = train_test_split(X,Y,test_size=0.2,random_state=42)
print('X_train is \n',X_train,'\n\nY_train is \n',y_train )
```

X_train is							
	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
60	2	84	0	0	0	0.0	
618	9	112	82	24	0	28.2	
346	1	139	46	19	83	28.7	
294	0	161	50	0	0	21.9	
231	6	134	80	37	370	46.2	
..	...	...	...	...	...	...	
71	5	139	64	35	140	28.6	
106	1	96	122	0	0	22.4	
270	10	101	86	37	0	45.6	
435	0	141	0	0	0	42.4	
102	0	125	96	0	0	22.5	
DiabetesPedigreeFunction Age							
60		0.304	21				
618		1.282	50				
346		0.654	22				
294		0.254	65				
231		0.238	46				
..		...	...				

```
[614 rows x 8 columns]
```

```
60      0
618     1
346     0
294     0
231     1
      ..
71      0
106     0
270     1
435     1
102     0
Name: Outcome, Length: 614, dtype: int64
```

```
y_pred = dtc.predict(X_test)
y_pred
```

```
print("Accuracy: ",round(metrics.accuracy_score(y_test,y_pred)*100,2),"%")
```

→ Accuracy: 75.97 %

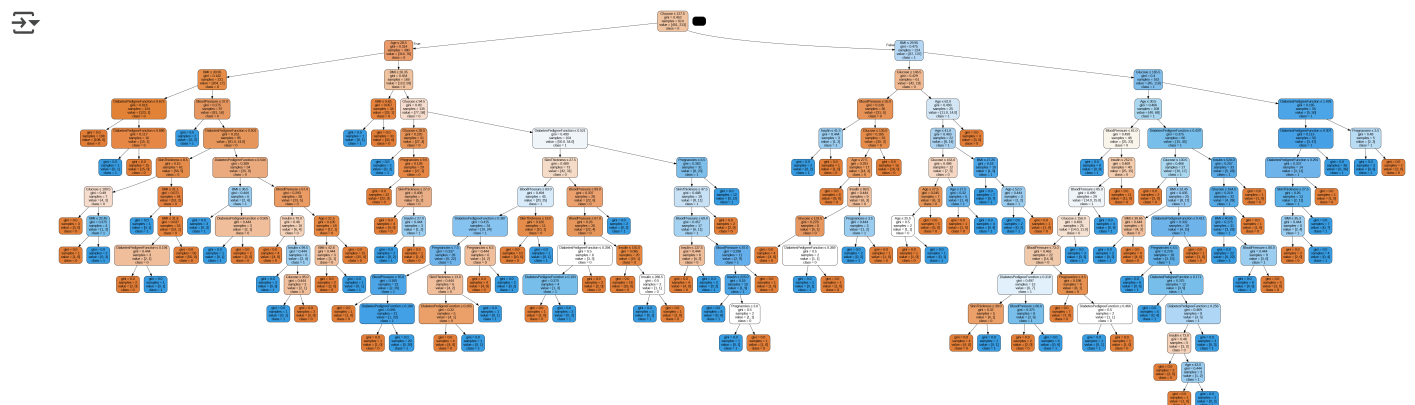
▼

```

[+] Requirement already satisfied: pydotplus in /usr/local/lib/python3.11/dist-packages (2.0.2)
Requirement already satisfied: pyparsing>=2.0.1 in /usr/local/lib/python3.11/dist-packages (from pydotplus) (3.2.1)

```

```
dot_data = StringIO()
export_graphviz(dtc, out_file=dot_data, filled=True, rounded=True, special_characters=True, feature_names=X.columns, class_names=['0', '1'])
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
graph.write_png('diabetes.png')
Image(graph.create_png())
```




### 3. Experiment with hyperparameters such as the maximum depth, minimum samples per leaf, and splitting criteria (e.g., Gini Index, Entropy).

```
dtc = DecisionTreeClassifier(criterion='entropy',max_depth=3)
dtc = dtc.fit(X_train,y_train)
```

```
y_pred = dtc.predict(X_test)
y_pred
```

```
array([0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0,
       0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0,
       0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1,
       0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0,
       0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0,
       0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1,
       0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0])
```

```
import seaborn as sns
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
# Confusion Matrix
plt.figure(figsize=(6,4))
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt="d", cmap="Blues", xticklabels=["No Diabetes", "Diabetes"], yticklabels=['No Diabetes', 'Diabetes'])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
# Model Evaluation
y_pred = dtc.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
```



```
Accuracy: 0.7662337662337663
Classification Report:
              precision    recall  f1-score   support

     0       0.81         0.84         0.82         99
     1       0.69         0.64         0.66         55

 accuracy          0.77         0.77         0.76         154
 macro avg         0.75         0.74         0.74         154
 weighted avg         0.76         0.77         0.76         154
```

```
import matplotlib.pyplot as plt
from sklearn.tree import plot_tree
# Visualizing the Decision Tree
plt.figure(figsize=(12, 8))
plot_tree(dtc, feature_names=X.columns, class_names=["No Diabetes", "Diabetes"], filled=True)
plt.show()
```

